Russell Dudek

AI Operations & Workflow Systems Leader

412.287.8640 · russelldudek@gmail.com
linkedin.com/in/russelldudek · Pittsburgh, Pennsylvania
https://russelldudek.github.io/1Mind/

First 90 Days

1mind · AI Operations Lead
Earn the right to scale

Entry thesis: Build trust and leverage in the same motion. Begin with observed Sales work, qualify a portfolio rather than chasing anecdotes, run two bounded learning cycles, and leave behind an operating system that makes the next workflow faster, safer, and easier to own.

OPERATING PRINCIPLES

Work before tools
Start with the decision, handoff, owner, baseline, and failure state.

Two speeds
Pair a quick learning cycle with one integration-dependent pattern.

Authority explicit
Name what AI may do, what humans retain, and who owns production.

Evidence over theater
Measure retained use, workflow outcome, exceptions, economics, and reuse.

WHAT I MUST LEARN BEFORE PRESCRIBING

Commercial work
Where Sales time, buyer continuity, follow-up, demos, handoffs, proposals, or onboarding currently lose context or create avoidable rework.

Existing capability
Which internal AI tools, automations, experiments, APIs, data sources, and product capabilities already exist—and what people trust or bypass.

Decision rights
What the AI Operations Lead owns directly; what remains with RevOps, Engineering, Security, Legal, Finance, and workflow owners.

Scale economics
Where maintenance, model/tool cost, integration effort, support demand, and change burden could erase an apparent quick win.

WEEKS 1–4 · READ, ALIGN, AND BOUND

WEEK 1
Align sponsorship, role charter, decision rights, operating cadence, and access.
Output: stakeholder and authority map

WEEK 2
Observe representative Sales workflows, queues, handoffs, current AI use, and workarounds.
Output: operating-surface map

WEEK 3
Baseline candidate workflows and populate value, context, risk, adoption, ownership, and reuse fields.
Output: ranked opportunity register

WEEK 4
Select two workflows and agree on human authority, evaluation, stop criteria, and Engineering handoff.
Output: signed workflow charters

DAY-30 DECISION GATE

Do we have two consequential, observable workflows with committed owners, explicit authority, usable context, and a credible path to evidence?

If not, do not rush into a prototype. Repair the prerequisite or replace the workflow.

Run two learning cycles

Days 31–60
Build + prove

PILOT PORTFOLIO DESIGN

PILOT A · QUICK LEARNING CYCLE

Low integration, fast behavioral evidence

Use an existing platform or lightweight automation to test context quality, human review, workflow fit, adoption, and value without pretending the result is production infrastructure.

- Short path to live work
- High learning value
- Clear owner and rollback

PILOT B · INTEGRATION-DEPENDENT PATTERN

Higher leverage, explicit Engineering boundary

Test a workflow whose durable value depends on systems integration, access, reliability, observability, or production ownership—so the handoff model is proven early.

- Architecture and security review
- Nonfunctional requirements
- Production owner or explicit hold

DAYS 31–60 WORKSTREAMS

Workstream	Actions	Evidence / output
Context contract	Define required inputs, missing-data behavior, source freshness, access, privacy, and how the intervention exposes uncertainty rather than silently filling gaps.	Context map and missing-data rules
Human authority	Specify what AI may retrieve, summarize, draft, recommend, classify, or execute; name approvals, commitments, exceptions, and escalation that remain human-owned.	Authority and exception matrix
Bounded build	Prototype the smallest useful intervention with visible failure states, logging, rollback, and a documented total operating-cost hypothesis.	Two usable prototypes
Live-work pilot	Place each intervention in the existing Sales cadence with a small representative cohort; do not create a separate destination unless the workflow requires it.	Usage, correction, exception, and outcome evidence
Adoption system	Train in context, create short standard work, establish feedback and support, and make the workflow owner—not AI Operations—the long-term behavior owner.	Owner-ready adoption package
Engineering gate	Translate the mechanism into architecture, integration, security, reliability, observability, release, and support requirements where custom development is warranted.	Hardening brief or explicit hold decision

WEEKLY PILOT CADENCE

Monday
Workflow signal, backlog, owner decisions

Midweek
User observation, corrections, exceptions, support burden

Friday
Outcome evidence, learning, next experiment or rollback

Biweekly gate
RevOps + Engineering decision on integration and ownership

DAY-60 DECISION GATE

Did the intervention change work in a way users retain—and is the durable value worth the control, maintenance, and integration burden?

Possible decisions: continue, integrate, revise, narrow, hold, or stop. A fast no is a learning asset when the evidence is preserved.

Standardize what earned it

Days 61–90
Compound + transfer

DAYS 61–90 WORKSTREAMS

Workstream	Actions	Evidence / output
Disposition	Standardize or integrate the pattern that earned it. Revise, narrow, hold, or stop the other based on evidence—not sponsorship volume or prototype novelty.	Decision record with rationale and owner
Operating memory	Version reusable context schemas, prompts, components, guardrails, evaluation, exception patterns, training, ownership, support, and measurement assets.	Reusable asset package and change record
Ownership transfer	Move normal operation, support, reinforcement, and outcome accountability to the workflow and technical owners; define AI Operations' ongoing role.	Runbook, support boundary, and ownership acceptance
Portfolio governance	Launch a monthly AI Operating Memory Review for opportunity intake, release disposition, adoption, exceptions, economics, Engineering demand, and learning.	Visible portfolio and review cadence
GTM expansion	Select the next adjacent workflow based on evidence and reuse potential; identify which parts transfer and which require fresh discovery.	Next-quarter portfolio and capacity decision

90-DAY OPERATING SCORECARD

Dimension	Evidence expected by day 90	Decision enabled
Business outcome	A workflow-specific baseline and observed change in cycle time, first-pass quality, rework, continuity, or another relevant unit.	Is the work worth continuing?
Retained adoption	Repeated use in the real workflow, named owner, support pattern, and user-friction record.	Did the intervention become a habit?
Trust + control	Authority, overrides, missing context, exceptions, escalation, rollback, and incident expectations are visible.	Can people use it confidently?
Economics	Tool/model cost, labor avoided, maintenance, support, rework, and Engineering demand are considered together.	Is durable value worth total cost?
Reuse + learning	Context, prompts, components, guardrails, training, and evidence improve the next workflow.	Did the company become smarter?

PRIMARY RISKS AND COUNTERMEASURES

Shadow automation	Make ownership, data boundaries, failure states, support, and production handoff explicit before scale.
Prototype theater	Require live-work evidence, retained use, workflow outcome, and operating-cost visibility.
Change fatigue	Choose consequential work, minimize new destinations, train in context, and stop low-value pilots quickly.
Engineering overload	Qualify integration demand early; provide bounded requirements and capacity decisions instead of vague asks.
False transfer	Reuse assets, not assumptions; re-observe each workflow before standardizing across teams.

DAY-90 DELIVERABLES

Role charter + decision rights	**Sales workflow and context maps**	**Ranked opportunity portfolio**
Two pilot decision records	**Reusable operating-memory asset package**	**Monthly review and next-quarter plan**

These are proposed operating outputs and discovery hypotheses, not guarantees or claims about 1mind's undisclosed baseline.